Assignment 15
LLMOps, Deployment & Safety

Part C — Building GenAI Applications

9 questions

★ **Must-know** — asked very frequently in GenAI / ML interviews (3 questions). Prepare these first.

1. How do you deploy an LLM application to production? Describe the architecture. ★
2. How do you reduce latency and cost (caching, smaller models, streaming, batching, prompt compression, routing)? ★
3. What are vLLM, TGI and Ollama? What is continuous batching and PagedAttention?
4. What is Groq / LPU inference and why is its latency so low?
5. What is semantic caching?
6. What are guardrails? How do you handle PII, toxic content and jailbreaks? ★
7. How do you monitor an LLM app in production (latency, tokens, cost, feedback, drift)?
8. What is Gradio / Streamlit used for? Build a quick demo UI for your RAG app.
9. Which speech-to-text and text-to-speech models have you used?